

Statistical Modeling

CH.1 - Simple Linear Regression

SS 2022 || Prof. Dr. Buchwitz

Wir geben Impulse

1 Organizational Information

2 Introduction

3 Simple Linear Regression

Lecturer

Professor Benjamin Buchwitz

- Room 2.3.14, Lindenstr. 53, Meschede
- Email: Buchwitz.Benjamin@fh-swf.de

Unit objectives

- 1 To obtain an understanding of common statistical methods used in statistical modeling.
- 2 To develop the computer skills required to model relationships found in business, economic and social sciences contexts;
- 3 To gain insights into the problems of implementing and conducting analyses for professional use.

Course Contents

Session	Topic
1	Simple Linear Regression
2	Multiple Linear Regression
3	Regression Diagnostics
4	Qualitative Variables as Predictors
5	Transformation of Variables
6	Weighted Least Squares
7	Correlated Errors
8	Analysis of Collinear Data
9	Working with Collinear Data
10	Variable Selection Procedures
11	Logistic Regression
12	Further Topics



Install R

<https://cloud.r-project.org/>

Install RStudio

<https://www.rstudio.com/products/rstudio/download/#download>

Grading is based on a portfolio examination with three parts:

- 1 One Lecture Recap Presentation (20%)
- 2 Hand-in Exercises (40%)
- 3 Final Case Study (40%)

1 Organizational Information

2 Introduction

3 Simple Linear Regression

What is Regression Analysis?

- Regression analysis is a conceptually simple method for investigating functional relationships among variables.
- The relationship is expressed in the form of an equation or a model connecting the **response** or **dependent variable** with one or more **explanatory** or **predictor** variables.
- We denote the response variable by Y and the set of predictor variables by X_1, X_2, \dots, X_p , where p denotes the number of predictor variables.
- The **true** relationship between the response and its predictors can be approximated by the regression model, where ϵ represents the random discrepancy in the relation.

$$Y = f(X_1, X_2, \dots, X_p) + \epsilon$$

The Regression Formula

- The function $f(X_1, X_2, \dots, X_p)$ describes the relationship between Y and X_1, X_2, \dots, X_p and can take any functional form.
- One example of a function is the linear regression model:

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots \beta_p X_p + \epsilon$$

- Here $\beta_0, \beta_1, \dots, \beta_p$ are called the regression parameters or coefficients, which are unknown constants and need to be estimated from data.

Data Example: River Data

- Nitrogen: Mean nitrogen concentration (mg/l) based on samples taken at regular intervals during the spring, summer and fall months
- Agr: Percentage of land area currently in agricultural use
- Forest: Percentage of forest land
- Rsdntial: Percentage of land area in residential use
- ComIndl: Percentage of land area in either commercial or industrial use

```
head(P010)
```

##	Agr	Forest	Rsdntial	ComIndl	Nitrogen
## Olean	26	63	1.2	0.29	1.10
## Cassadaga	29	57	0.7	0.09	1.01
## Oatka	54	26	1.8	0.58	1.90
## Neversink	2	84	1.9	1.98	1.00
## Hackensack	3	27	29.4	3.11	1.99
## Wappinger	19	61	3.4	0.56	1.42

Data Example: Motor Trend US Car Magazine

see help(mtcars) for variable description

mtcars

##	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
## Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
## Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
## Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
## Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
## Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
## Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
## Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
## Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
## Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
## Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
## Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
## Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
## Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
## Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
## Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
## Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
## Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
## Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
## Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
## Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
## Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
## Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2

Steps in Regression Analysis

- 1 Statement of the problem
- 2 Selection of potentially relevant variables
- 3 Data collection
- 4 Model specification
- 5 Choice of fitting method
- 6 Model fitting
- 7 Model validation and criticism
- 8 Using the chosen model(s) for the solution of the proposed problem

Statement of the problem

- Every analysis starts with the definition of the problem, which includes formulation of questions addressed by the analysis.
- Ill-defined problems or misformulated questions can lead to wasted effort or the selection of a wrong model.
- Finding and formulating suitable questions is probably the hardest part in an analysis.

Example: Problem Statement Definition

- Assume we want to research whether or not an employer is discriminating against a group of employees, e.g. women and data on salary, gender and qualification is available.
- There are multiple definitions of discriminations available in the literature (a) women are paid less than equally qualified men, or (b) women are more qualified than equally paid men.

Example: Problem Statement Definition

- Assume we want to research whether or not an employer is discriminating against a group of employees, e.g. women and data on salary, gender and qualification is available.
- There are multiple definitions of discriminations available in the literature (a) women are paid less than equally qualified men, or (b) women are more qualified than equally paid men.

Your turn

What is the modeling implication of the definition?

Example: Problem Statement Definition

- Assume we want to research whether or not an employer is discriminating against a group of employees, e.g. women and data on salary, gender and qualification is available.
- There are multiple definitions of discriminations available in the literature (a) women are paid less than equally qualified men, or (b) women are more qualified than equally paid men.


Your turn

What is the modeling implication of the definition?

a) $salary = f(qualification, gender) + \epsilon$

b) $qualification = f(salary, gender) + \epsilon$

Flowchart



1 Organizational Information


2 Introduction

3 Simple Linear Regression

$$Y = f(X) + \epsilon$$

- We start with the simple case to study the relationship between the response Y and a single predictor X .
- As we only have one regressor variable we drop the subscript to simplify the notation ($X_1 = X$).
- We derive and formulate the regression model and focus on the key results but favor numerical examples over mathematical derivations.

Covariance



Determine the sign:

- $y_i - \bar{y}$ the deviation of each observation y_i from the mean of the response variable,
- $x_i - \bar{x}$ the deviation of each observation x_i from the mean of the predictor variable, and
- the product of the above quantities, $(y_i - \bar{y})(x_i - \bar{x})$

Quadrant	$y_i - \bar{y}$	$x_i - \bar{x}$	$(y_i - \bar{y})(x_i - \bar{x})$
1 (top right)			
2 (top left)			
3 (bottom left)			
4 (bottom right)			

Positive Relationship

- If the linear relationship between Y and X is **positive** (when X increases, Y also increases), then there are more points in the first and third quadrants than in the second and fourth.
- The sum over the elements in the last column is likely to be positive, that is $\text{Cov}(Y, X) > 0$.

Positive Relationship

- If the linear relationship between Y and X is **positive** (when X increases, Y also increases), then there are more points in the first and third quadrants than in the second and fourth.
- The sum over the elements in the last column is likely to be positive, that is $\text{Cov}(Y, X) > 0$.

Negative Relationship

- If the linear relationship between Y and X is **negative** (as X increases Y decreases), then there are more points in the second and fourth quadrants than in the first and third.
- The sum over the elements in the last column is likely to be negative, that is $\text{Cov}(Y, X) < 0$.

$$\text{Cov}(X, Y) = \frac{1}{n-1} \sum_{i=1}^n (y_i - \bar{y})(x_i - \bar{x})$$

- The quantity calculated using the above formula is called the covariance.
- The sign of the covariance indicates the relationship between Y and X.
- The covariance can **only indicate the direction** of a relationship, and does not tell much about the strength of the relationship.
- the covariance is unit sensitive, changing the unit of a measurement (e.g. from Euro to kEuro) changes the value of the covariance.

Your turn

What happens if we calculate $\text{Cov}(X, Y)$ instead of $\text{Cov}(Y, X)$?

- To avoid the obvious disadvantages of the covariance we can standardize (z-transform) each variable before computing the covariance.
- Standardizing Y means subtracting the mean \bar{y} and dividing by the associated sample standard deviation s_y .
- The resulting variable z_i has mean zero and unit standard deviation.

$$z_i = \frac{y_i - \bar{y}}{s_y} \quad \text{with} \quad s_y = \sqrt{\frac{\sum_{i=1}^n (y_i - \bar{y})^2}{n - 1}}$$

$$\text{Cor}(Y, X) = \frac{1}{n-1} \sum_{i=1}^n \left(\frac{y_i - \bar{y}}{s_y} \right) \left(\frac{x_i - \bar{x}}{s_x} \right) = \frac{\text{Cov}(Y, X)}{s_y s_x}$$

- Calculating the covariance of the standardized values yields the correlation coefficient.
- $\text{Cor}(Y, X)$ can be interpreted in two ways, either as
 - ▶ the covariance between two standardized variables or as
 - ▶ the ratio of the covariance to the standard deviations of the two variables
- Opposed to the covariance, $\text{Cor}(Y, X)$ is scale invariant so that it is not affected by unit changes. It also satisfies $-1 \geq \text{Cor}(Y, X) \geq 1$ and therefore indicates **direction** and **strength**.


Correlation Coefficient

$Cor(Y, X) = 0$ does not necessarily mean that the variables are not related!

```
x <- seq(from=-5, to=5, by=.1)
y <- 50 - x^2
cor_yx = cor(y,x)
round(cor_yx, digits=4)
```

```
## [1] 0
```

```
plot(x,y)
```



Example: Anscombe Quartet

```
knitr::kable(anscombe[,c("y1", "x1", "y2", "x2", "y3", "x3", "y4", "x4")], booktabs=T)
```


y1	x1	y2	x2	y3	x3	y4	x4
8.04	10	9.14	10	7.46	10	6.58	8
6.95	8	8.14	8	6.77	8	5.76	8
7.58	13	8.74	13	12.74	13	7.71	8
8.81	9	8.77	9	7.11	9	8.84	8
8.33	11	9.26	11	7.81	11	8.47	8
9.96	14	8.10	14	8.84	14	7.04	8
7.24	6	6.13	6	6.08	6	5.25	8
4.26	4	3.10	4	5.39	4	12.50	19
10.84	12	9.13	12	8.15	12	5.56	8
4.82	7	7.26	7	6.42	7	7.91	8
5.68	5	4.74	5	5.73	5	6.89	8

Your turn


Choose one of the datasets (e.g. $i = 3$) and calculate \bar{y}_i , \bar{x}_i , $\text{Cov}(y_i, x_i)$ and $\text{Cor}(y_i, x_i)$ using R (**Hint:** `mean`, `cov`, `cor`).

Results: Anscombe Quartet


(a)




(b)



(c)



(d)



- Like many other summary statistics the correlation coefficient can be substantially influenced by one of a few outliers in the data.
- All four datasets in the Anscombe quartet have almost the **same summary** statistics, despite being inherently different.
- A purely descriptive analysis can not reveal the different patterns **we need to plot the data before starting an analysis.**
- Findings:
 - a) can be adequately described by a linear model
 - b) is nonlinear and would be better fitted by a quadratic function
 - c) one outlier distorts the slope and intercept of the lines
 - d) is unsuitable for linear fitting as the line is determined by a single extreme observation

Example: Computer Repair Data

```
# Minutes = Duration of the service operation  
# Units = Number of computers repaired during service operation  
head(P031)
```

```
##   Minutes Units  
## 1      23     1  
## 2      29     2  
## 3      49     3  
## 4      64     4  
## 5      74     4  
## 6      87     5
```

Your turn

Calculate $\text{Cov}(Y, X)$ and $\text{Cor}(Y, X)$ manually (step-by-step) using R by avoiding the internal functions `cov` and `cor`.

The Simple Linear Regression Model

- The correlation coefficient is useful to measure the strength of a pairwise relationship, it **cannot be used for prediction purposes**.
- That means that we cannot use $Cor(Y, X)$ to predict one variable, when the other one is given.
- Regression is an extension to correlation analysis and cannot only measure direction, but allows for **numerically describing** that relationship.

The Simple Linear Regression Model

$$Y = \beta_0 + \beta_1 X + \epsilon$$

- β_0 and β_1 are constants called the regression coefficients, and ϵ is the error term.
 - ▶ β_0 is called the intercept. It is the prediction value, when $X = 0$.
 - ▶ β_1 is called the slope. It can be interpreted as the change in Y , when X changes by one unit.
- Each observation in the data can therefore be written as:

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i \quad \text{with} \quad i = 1, 2, \dots, n$$

The Simple Linear Regression Model

- We assume that (in the range of our observations studied), the linear equation provides an **acceptable approximation** to the real relationship: Y is approximately a linear function of X .
- The error term ϵ measures the discrepancy of the approximation.
- That simple linear regression model is linear in two ways:
 - ▶ the relationship between X and Y is linear
 - ▶ more generally the word linear describes that the regression parameters β_0 and β_1 enter the equation in a linear fashion
 - ▶ $Y = \beta_0 + \beta_1 X^2 + \epsilon$ is still a linear model but with a quadratic term!
- In correlation X and Y are of equal “importance” which is reflected in the symmetry $\text{Cor}(Y, X) = \text{Cor}(X, Y)$.
- In regression we want to explain Y , hence the importance of the predictor X lies on its ability to account for the variability of the response.

Example: Computer Repair Data

Reconsidering the computer repair data and assuming we want to predict the numbers of support engineers that will be required for a task, we can now formulate an equation in form of a linear model that is assumed to represent the relationship between the length of service calls and the number of electronic components in the computer that must be repaired.

$$\text{Minutes} = \beta_0 + \beta_1 \text{ Units} + \epsilon$$


Your turn

Consult the scatter plot (plot) of the data (P031) and check whether the straight linear relationship is a reasonable assumption.

How do we
determine β_0 and β_1 ?

Parameter Estimation

```
plot(P031$Minutes,P031$Units,xlab="Minutes",ylab="Units", pch=19)
```



- We want values for β_0 and β_1 that give the *best fit* or the *best representation* for the points in the graph.
- This can be achieved using the **least squares method** that minimizes the sum of squares of **vertical distances**.
- Those vertical distances from each point to the line represent the errors ϵ_i and can be obtained by:

$$\epsilon_i = y_i - \beta_0 - \beta_1 x_i \quad \text{for } i = 1, 2, \dots, n$$

Parameter Estimation

- As β_0 and β_1 are unknown, but required to calculate the errors and therefore the sum of squared errors, we can devise a function for that:

$$S(\beta_0, \beta_1) = \sum_{i=1}^n \epsilon_i^2 = \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2$$

- This is a quadratic function that can be minimized. The analytical solution for the values $\hat{\beta}_0$ and $\hat{\beta}_1$ that minimize the function $S()$ are

$$\hat{\beta}_1 = \frac{\sum (y_i - \bar{y})(x_i - \bar{x})}{\sum (x_i - \bar{x})^2} \quad \text{and} \quad \hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$$

- Both, $\hat{\beta}_0$ and $\hat{\beta}_1$ are called the **least squares estimates** and give the line with the smallest possible sum of squares of vertical distances.

- The **least squares regression line** can always be found (does always exist) and is given by

$$\hat{Y} = \hat{\beta}_0 + \hat{\beta}_1 X$$

- For each observation we can compute a **fitted value**, which is given by

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i \quad \text{for } i = 1, 2, \dots, n$$


- Each point (x_i, \hat{y}_i) is a point **on the regression line**
- The corresponding vertical distances are called **ordinary least squares residuals** and can be computed like

$$\hat{\epsilon}_i = y_i - \hat{y}_i \quad \text{for } i = 1, 2, \dots, n$$


Parameter Estimation

Your turn

Add a sketch of the least squares regression line to the plot above and include, mark and annotate the fitted values and the ordinary least squares residuals.



Parameter Estimation



Your turn

- Calculate $\hat{\beta}_0$ and $\hat{\beta}_1$ twice using R.
 - 1) Manually (abstain from `cor` and `cov`) using R
 - 2) Using the functions mentioned above
- Plot the data and add your calculated regression line to that plot (**Hint:** `abline`)

- So far we only made one assumption or hypothesis about the relationship between the response and predictor variables, which is called the **linearity assumption**.
- An early step in an analysis should always be the validation of this assumption: *We wish to determine if the data at hand supports the assumption that Y and X are linearly related.*
- An **informal** way to check this assumption is to check the scatter plot.
- A more **formal** way to check the assumption and to measure the usefulness of X as a predictor for Y is to conduct a hypothesis test about the regression parameter β_1 .

- Testing for the postulated relationship can be done by checking the hypothesis that $\beta_1 = 0$, which means that there is **no linear relationship** between X and Y .
- Finding that $\beta_1 > 0$ or $\beta_1 < 0$ is equivalent to $\beta_1 \neq 0$ and would provide **evidence (not proof!) for an existing linear relationship**.
- Testing of this hypothesis requires the assumption that the errors ϵ_i are independent random quantities originating from a normal distribution with zero mean and common variance σ^2 .
 - ▶ $\epsilon \sim N(0, \sigma^2)$
 - ▶ ϵ_i are independent

- Given that the assumptions for the error term ϵ hold, $\hat{\beta}_0$ and $\hat{\beta}_1$ are unbiased estimates of β_0 and β_1 .
- This means that $\hat{\beta}_0$ and $\hat{\beta}_1$ allow to draw conclusions about the unobserved and unknown parameters β_0 and β_1 in the population, hence $E(\hat{\beta}) = \beta$.
- Under the mentioned circumstances the variances of the regression coefficients are

$$\text{Var}(\hat{\beta}_0) = \sigma^2 \left[\frac{1}{n} + \frac{\bar{x}^2}{\sum (x_i - \bar{x})^2} \right] \quad \text{and} \quad \text{Var}(\hat{\beta}_1) = \frac{\sigma^2}{\sum (x_i - \bar{x})^2}$$

- The variances of $\hat{\beta}_0$ and $\hat{\beta}_1$ depend on the unknown and unobservable parameter σ^2 , which needs to be estimated from the data before we can proceed.

- An unbiased estimate of σ^2 is given by

$$\hat{\sigma}^2 = \frac{\sum \epsilon_i^2}{n - 2} = \frac{\sum (y_i - \hat{y}_i)^2}{n - 2} = \frac{SSE}{n - 2}$$

- Here *SSE* is an abbreviation for Sum of Squares Error (Residuals).
- The number $n - 2$ is called *degrees of freedom (df)* and is equal to the number of observations n minus the number of estimated regression coefficients.

- Plugging $\hat{\sigma}^2$ into $\text{Var}(\hat{\beta}_0)$ and $\text{Var}(\hat{\beta}_1)$ yields unbiased estimates of the respective variances.
- The estimate of the standard deviation is called the **standard error (s.e.)**

$$s.e.(\hat{\beta}_0) = \hat{\sigma}^2 \sqrt{\frac{1}{n} + \frac{\bar{x}^2}{\sum (x_i - \bar{x})^2}} \quad \text{and} \quad s.e.(\hat{\beta}_1) = \frac{\hat{\sigma}^2}{\sqrt{\sum (x_i - \bar{x})^2}}$$

- The standard error of $\hat{\beta}_1$ is a measure of how precisely the slope has been estimated. The smaller $s.e.(\hat{\beta}_1)$, the more precise is the estimator.

We are now in the position to perform statistical analysis concerning the usefulness of X as a predictor of Y . Under the assumption of normality, an appropriate test for testing the hypothesis is the t-test.

$$H_0 : \beta_1 = 0 \quad \text{versus} \quad H_1 : \beta_1 \neq 0$$

The test statistic follows a Student t distribution with $n - 2$ degrees of freedom and we need a specified significance value (e.g. $\alpha = 0.05$) to perform the test.

$$t_1 = \frac{\hat{\beta}_1}{s.e.(\hat{\beta}_1)}$$


Carrying out the test is done by comparing the value t_1 against the appropriate critical value obtained from the t-table, which is $t_{(n-2, \alpha/2)}$ (Note that we divide α by 2 as we have a two-sided test).

Reject H_0 at the given significance level if:

$$t_1 \geq t_{(n-2, \alpha/2)} \quad \text{or} \quad t_1 \leq -t_{(n-2, \alpha/2)}$$

One condition is fulfilled if $|t_1| \leq t_{(n-2, \alpha/2)}$. A criterion equivalent to that is to compare the p-value (implicit probability value) for the t-test with α and reject H_0 if $p(|t_1|) \leq \alpha$, where $p(|t_1|)$, called the p-value, is the sum of the two shaded areas under the following curve. This value is also provided by R.

Tests of Hypotheses



The t-test can be generalized to test the more general hypothesis $H_0 : \beta_1 = \beta_1^0$, where β_1^0 is a constant chosen by the data analyst.

$$t_1 = \frac{\hat{\beta}_1 - \beta_1^0}{s.e.(\hat{\beta}_1)}$$

The t-test can also be used for testing the intercept β_0 in the same fashion.

Tests of Hypotheses

Your turn

Identify all quantities in the showed regression output.

```
summary(lm(Minutes ~ 1 + Units, data=P031))

##
## Call:
## lm(formula = Minutes ~ 1 + Units, data = P031)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9.232 -3.341 -0.714  4.777  7.803
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    4.162      3.355    1.24    0.24
## Units         15.509      0.505   30.71 8.9e-13 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5.39 on 12 degrees of freedom
## Multiple R-squared:  0.987, Adjusted R-squared:  0.986
## F-statistic: 943 on 1 and 12 DF, p-value: 8.92e-13
```

Based on the assumption that the errors ϵ are normally distributed $\epsilon \sim N(0, \sigma^2)$, we concluded that the sampling distribution of β_0 and β_1 is also normal.

- The $(1 - \alpha) \cdot 100\%$ confidence intervals for β_0 and β_1 are given by

$$\hat{\beta}_0 \pm t_{(n-2, \alpha/2)} \cdot \text{s.e.}(\hat{\beta}_0) \quad \text{and} \quad \hat{\beta}_1 \pm t_{(n-2, \alpha/2)} \cdot \text{s.e.}(\hat{\beta}_1)$$

- ▶ with $t_{(n-2, \alpha/2)}$ being the $(1 - \alpha/2)$ percentile of a t distribution with $n - 2$ degrees of freedom
- ▶ **Interpretation:** If we took repeated samples of the same size (n) at the same values of X and construct e.g. the 95% confidence intervals for the slope parameter (based on $\hat{\beta}_1$) for each sample, then 95% of these intervals would be expected to contain the true but unknown value β_1 of the slope.

Confidence Intervals

```
summary(lm(Minutes ~ 1 + Units, data=P031))
```

```
##
## Call:
## lm(formula = Minutes ~ 1 + Units, data = P031)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9.232 -3.341 -0.714  4.777  7.803
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    4.162      3.355     1.24    0.24
## Units         15.509      0.505    30.71 8.9e-13 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5.39 on 12 degrees of freedom
## Multiple R-squared:  0.987, Adjusted R-squared:  0.986
## F-statistic: 943 on 1 and 12 DF, p-value: 8.92e-13
```

Your turn

Construct the 95% confidence intervals for the regression parameters and interpret their values.

In addition to describing and explaining observed relationships, the fitted regression equation can be used for prediction. We distinguish two types of predictions:

- 1) The prediction of the **value** of the response Variable Y which corresponds to any chosen value x_0 of the predictor variable.
- 2) The estimation of the **mean** response μ_0 , when $X = x_0$.

Value of Response & Prediction Limits

The predicted value \hat{y}_0 is given by $\hat{y}_0 = \hat{\beta}_0 + \hat{\beta}_1 x_0$. Its confidence interval can be constructed by $\hat{y}_0 \pm t_{(n-2, \alpha/2)} \cdot s.e.(\hat{y}_0)$. The standard error of the prediction is

$$s.e.(\hat{y}_0) = \hat{\sigma} \sqrt{1 + \frac{1}{n} + \frac{(x_0 - \bar{x})^2}{\sum (x_i - \bar{x})^2}}$$

Mean Response & Confidence Limits

The predicted value $\hat{\mu}_0$ is given by $\hat{\mu}_0 = \hat{\beta}_0 + \hat{\beta}_1 x_0$. Its confidence interval is given by $\hat{\mu}_0 \pm t_{(n-2, \alpha/2)} \cdot s.e.(\hat{\mu}_0)$. The standard error of the prediction is

$$s.e.(\hat{\mu}_0) = \hat{\sigma} \sqrt{\frac{1}{n} + \frac{(x_0 - \bar{x})^2}{\sum (x_i - \bar{x})^2}}$$

- How do those two predictions differ from each other?

- How do those two predictions differ from each other?
- What do they have in common?

- How do those two predictions differ from each other?
- What do they have in common?
- Why is the uncertainty smaller when predicting the mean response $\hat{\mu}_0$ instead of the single value \hat{y}_0 ?

Your turn

Use the computer repair data (P031) to do the following with:

- Predict the length of a service call and the associated standard deviation in which four components have to be repaired
- Estimate the expected mean service time for calls that needed four components to be repaired.

Your turn

Use the computer repair data (P031) to do the following with:

- Predict the length of a service call and the associated standard deviation in which four components have to be repaired
- Estimate the expected mean service time for calls that needed four components to be repaired.


There are two dangers in such calculations:

- Notice the substantial uncertainty / the large standard error
- The linear relationship may only hold in the range of the observed data. In case of the computer repair data we would for example not use this procedure to predict service times for services which require e.g. 25 components to be repaired


Goodness of Fit

Which line has a higher quality of fit and therefore resembles the underlying relationship better?

(a)



(b)



The quality of the fit can be assessed by one of the following (highly related) ways:

- 1 The discussed tests (if H_0 is rejected), the magnitude of the values gives us information about the strength (not only existence) of the linear relationship between Y and X . The larger $|t|$ (or the smaller the p -value), the stronger the linear relationship. The tests are objective but require the stated assumption of normality of ϵ .
- 2 One can also revert to $\text{Cor}(Y, X)$ and the scatter plot of the data. The **closer** the set of points to a straight line, the closer is $|\text{Cor}(Y, X)|$ to 1 and the stronger is the linear relationship between Y and X . This approach is informal and subjective, but requires only the linearity assumption.

- 3 Examine the scatter plot of the response Y versus the fitted values \hat{Y} . The closer the points to a straight line, the stronger is the linear relationship. This can be measured using the correlation $\text{Cor}(\hat{Y}, Y)$. In simple linear regression this is redundant, as $\text{Cor}(\hat{Y}, Y) = |\text{Cor}(Y, X)|$. However when multiple regressors are available this is an informative plot.
- 4 Decomposition of the variance $\text{Var}(Y)$, which is in fact closely related to the previous approach of using $\text{Cor}(\hat{Y}, Y)$, but useful for simple and multivariate linear regression.

After computing the least squares estimates let us compute these quantities.


$$SST = \sum (y_i - \bar{y})^2$$

$$SSR = \sum (\hat{y}_i - \bar{y})^2$$

$$SSE = \sum (y_i - \hat{y}_i)^2$$

- Sum of Squares Total (SST) is the total deviation of Y from its mean \bar{y} .
- Sum of Squares Regression (SSR) is the explained deviation, that is modeled by the regression line.
- Sum of Squares Error (SSE) is the total deviation of the residuals.

Goodness of Fit (Graphical Illustration)



$$SST = SSR + SSE$$

$$SST = SSR + SSE$$

$$y_i = \hat{y}_i + (y_i - \hat{y}_i)$$

$$\text{Observed} = \text{Fit} + \text{Deviation from Fit}$$

$$SST = SSR + SSE$$

$$y_i = \hat{y}_i + (y_i - \hat{y}_i)$$

$$\text{Observed} = \text{Fit} + \text{Deviation from Fit}$$

Subtracting \bar{y} yields:

$$y_i - \bar{y} = (\hat{y}_i - \bar{y}) + (y_i - \hat{y}_i)$$

$$\text{Deviation from mean} = \text{Deviation due to Fit} + \text{Residual}$$

The sum of squared deviations in Y can be decomposed accordingly!

- The first, SSR, measures the quality of X as a predictor of Y
- the second, SSE, measures the error in this prediction.
- Therefore the ratio $R^2 = SSR/SST$ is the proportion of the total variation in Y that is accounted for by the predictor X . It follows that:

$$R^2 = \frac{SSR}{SST} = 1 - \frac{SSE}{SST}$$

Your turn

Calculate the R^2 for the Computer repair data and show that $R^2 = [\text{Cor}(Y, X)]^2 = [\text{Cor}(Y, \hat{Y})]^2$ holds in the case of simple linear regression.

- Regression Line through the Origin
- Trivial Models